

Millions of Games per Hour

Ben Deane

5th June 2014

Outline

- 1 History & Problem Space
- 2 Abstracting Matching
- 3 Cooperative Matching
- 4 Queueing and Distribution
- 5 Competitive Matching

History & Problem Space

- 1 History & Problem Space
- 2 Abstracting Matching
- 3 Cooperative Matching
- 4 Queueing and Distribution
- 5 Competitive Matching

History

- Server lists don't cut it any more
- Player population is large
- Game type proliferation
- e-Sports expectations

The Problem Space

- Find good games quickly
- Deal with varying game types
- Deal with large player population
- Deal with small player population

The Problem Space

- Assign games to servers
 - spread load evenly
 - fill up servers
 - deal with overload scenarios
- Deal with community preferences

Cooperative vs Competitive

Cooperative

- Drop-in, drop out
- Social matches
- Game type matching
- Party-based play

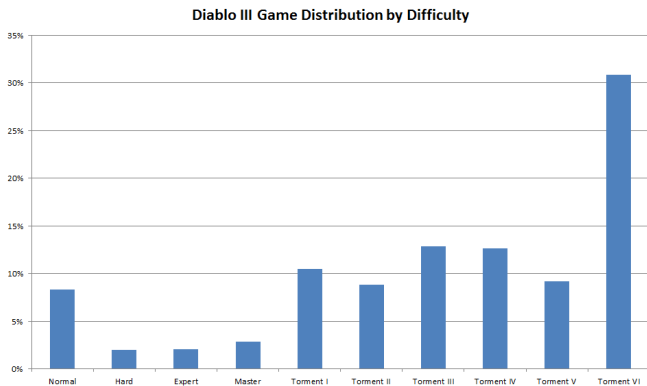
Competitive

- Join at start
- Skill-based matches
- Game size matching
- Team-based play

Evolution of players

- Players move through content over time
- Players gain skill over time
- Players return to content to farm it

Evolution of players




Battle.net Tenets

- Keep it simple
 - Functionality comes from composability, not monolithic behavior
 - The best code is no code
- Be reliable
 - Easy configuration
 - No single points of failure
- Be game agnostic

System Overview

Game clients

Game servers



Abstracting Matching

- 1 History & Problem Space
- 2 Abstracting Matching**
- 3 Cooperative Matching
- 4 Queueing and Distribution
- 5 Competitive Matching

What defines a game?

- A set of attributes

What defines a game?

- A set of attributes
- Partitioning attributes
 - Difficulty
 - Hardcore/Regular/Starter
 - Version


What defines a game?

- A set of attributes
- Partitioning attributes
 - Difficulty
 - Hardcore/Regular/Starter
 - Version
- Matchable attributes
 - Act number
 - Quest step
 - Other


Attributes

- Attributes are key-value pairs
- Battle.net doesn't know what they mean
- Battle.net knows how to
 - Wrangle them in data structures
 - Do computations with them (hashing, sorting, comparing)

What a client does




What a server does



Game Factories


- Game factories represent partitions
 - normal-nonhardcore-v201-factory
 - hard-nonhardcore-v201-factory
 - etc
- Game factories are
 - specified in configuration
 - instantiated in response to server connections
 - combinatorial on relatively few axes

When a server connects



I'll associate you with these factories!

When a client asks for a game



Game Factories

- Game factories reduce the matching problem
- Each factory matches the games it knows about
 - Based on the smaller number of matchable attributes
- Factories can use different strategies
- The factory abstraction is strategy-agnostic
 - cooperative
 - competitive


Cooperative Matching

- 1 History & Problem Space
- 2 Abstracting Matching
- 3 Cooperative Matching**
- 4 Queueing and Distribution
- 5 Competitive Matching

FIND a game

- The API deliberately says FIND a game
 - not join a game
 - not create a game
- The create/join dichotomy is not part of matchmaking
- If a game cannot be matched, one will be created
 - Either way, you get into a game
- CREATE and JOIN have their place, but it's not matchmaking

Matching on attributes



The most important "attributes"

- Is the game open for matching?
- Is there space in the game?
- Factories partition the open game list by number of open slots
- Players match in groups
 - individually
 - parties

Remaining problem

- We have a candidate set of games
 - that are open for matching
 - that can fit our players
 - that are associated with some attributes
- We want to match our attributes against the games


Dynamic matching

- N-dimensional nearest neighbor search?

Dynamic matching

- N-dimensional nearest neighbor search?
- Index the games list by each attribute
- Each (single) attribute lookup yields a set of games
- To find a match for all, compute the set intersection

Indexed games



Indexed stats too

- As for game matching, so for extracting stats
 - number of games
 - number of players
 - min/max/average game duration
 - etc
- Stats can be queried using the same attribute matching/indexing scheme

The Problems?

- Find good games quickly ✓
- Deal with varying game types ✓
- Deal with large player population ✓


The Problems?

- Find good games quickly ✓
- Deal with varying game types ✓
- Deal with large player population ✓
- Deal with small player population ?


The Problems?

- Find good games quickly ✓
- Deal with varying game types ✓
- Deal with large player population ✓
- Deal with small player population ?
- Assign games to servers

Filling servers



Spreading load



Spread players vs Fill up games

- Max N players in a game
- k players in a matching group
- Just match against games with the "right" number of open slots
 - to fill, match with $N-k, N-k-1, \dots, 1$
 - to spread, match with $1, 2, \dots, N-k$

Queueing and Distribution

- 1 History & Problem Space
- 2 Abstracting Matching
- 3 Cooperative Matching
- 4 Queueing and Distribution**
- 5 Competitive Matching


The Problems

- Take account of server load somehow
- Assign games to servers evenly
- Allow new servers to come online and get balanced
- Deal with servers being temporarily full


1st attempt

- One server deals with queueing
 - with fail over to another
- Poll game servers for load
- Assign a nominal load per game creation/player addition

1st attempt



Extra complexity



Extra complexity

- Game server capacity isn't the only factor
- Limited bandwidth to regional data centers results in poor experience
 - very important for hardcore mode!
 - KR-TW pipe is small

Problems with 1st attempt

- Polling has a delay
 - Queue master has to anticipate load assigned during the delay period
- Single point of failure
- Queue master doesn't know what the result of matching will be
 - it deals with groups of players only
 - it doesn't know whether a game will be joined or created
 - these two scenarios have different load characteristics

Problems with 1st attempt

- Hard to reason about rate of game influx to a given server
 - hard to bring up servers
 - open beta bug
- Hard to estimate load on servers
 - queueing is by player but load is by game
 - starting games is spiky load
 - running games is smooth load

2nd attempt

- Drop regional bandwidth requirement
- Allow servers to advertise the number of games they can take
- Slots apportioned to factories by popularity
- Queue can be distributed across hardware

Competitive Matching

- 1 History & Problem Space
- 2 Abstracting Matching
- 3 Cooperative Matching
- 4 Queueing and Distribution
- 5 Competitive Matching**

Competitive Matching

- Join at start
- Skill-based
 - Elo-like player rating
- 1v1 or NvN




Game agnosticism

- Battle.net doesn't know about player skill per se
 - Stats and logic are down to the game
 - Abstracted as a single player score

Hearthstone

- Batch players
- Sort batch by skill
- Make games according to threshold
- Easy!

Hearthstone

Batch	Sort	Form Games	
{phineas: 1.5} {ferb: 1.8}	{ candace: -0.7} {doof: -0.5}	{ candace: -0.7} {doof: -0.5}	
{isabella: 1.7} {doof: -0.5}	{vanessa: 0.3} {phineas: 1.5}	{vanessa: 0.3} {phineas: 1.5}	
{vanessa: 0.3} { candace: -0.7}	{isabella: 1.7} {ferb: 1.8}	{isabella: 1.7} {ferb: 1.8}	

StarCraft II

- More options make it harder
 - teams/random NvN/free-for-all
 - map selections
 - different player ping times
- Simple sorting doesn't work
 - hill-climbing optimizer

Hill Climbing

- 1 Take a batch of players, assemble some games
- 2 Swap something around
- 3 See if the games are better
- 4 If you still have time, goto 2
- 5 Start the games that are viable
- 6 Put leftovers in the next batch and relax constraints

Performance Issues

- Hill-Climbing algorithm is $O(n!)$
- Fewer players means you need to work harder
- More players means it's easier to make viable games
- With appropriate selections for batch size the system is self-regulating

Statistical Issues

- Player score starts out uncertain
 - 1v1 requires ~25 games to focus
 - Other modes require more games
- Unbalanced teams are hard to match
 - Teams of friends are often variable
 - e.g. Experienced player + novice
- Aggregating team scores is tricky

Statistical Issues

- Very good players can't get good matches
 - as in any sport
- Players like their rating to increase
- Players get better, then leave

Competitive Design Issues

- If your game is 1v1 and your matchmaker is perfect:
 - 50% of players will lose the first match
 - 25% of players will lose the first two matches
- You need to make the game fun even when players lose
 - Progression
 - Achievements

Abuse Issues

- Achievements incentivize loss-botting
- Disconnection = loss

Thanks for listening

- Factories allow abstraction of strategies
- Queueing to manage load
- Competitive and Cooperative are different animals
- Reduce problems by design

Ben Deane
bdeane@blizzard.com

Testing

- Bonus section!
- A little bit about how I tested. . .

The need to test

- I can't run at scale on my desktop machine
- I need to be sure that the system runs at scale

Possibilities

- Use repurposed ("spare") hardware
- Use Amazon EC2 or similar
- Or I could just figure out how to test on my machine

- I had already built the parts with unit tests
- I/O was separated out
- Configuration was dependency-injected
- Matchmaker logic was separated out

Unit testing for performance

- Absolute perf not so good
 - my machine isn't a production machine
 - my machine can't simulate a million players
 - unit tests are supposed to be fast

Unit testing for performance

- Modified unit test framework
 - Call tests N times and time the result
 - Vary N from 2^a to 2^b
 - Divide results to obtain complexity order
- Algorithmic complexity tests
 - No hidden $O(n)$ or worse algorithms
 - Everything is $O(\log_2 n)$
- I didn't do any statistical analysis: good enough is good enough

Thanks for listening (more)

- Factories allow abstraction of strategies
- Queueing to manage load
- Competitive and Cooperative are different animals
- Reduce problems by design
- Test at scale somehow

Ben Deane
bdeane@blizzard.com